

3rd Year Project Report

Catherine Vlasov

May 2018



Abstract

This is my abstract

Contents

1	Introduction	2
2	Background	3
2.1	Reinforcement Learning	3
2.2	Learning Algorithms	3
2.2.1	Monte Carlo Learning	3
3	Design & Implementation	4
3.1	Overview	4
3.2	Agents	4
3.3	Tic-Tac-Toe	5
3.3.1	Rules	5
3.3.2	Implementation Details	6
3.4	Chung Toi	7
3.4.1	Rules	7
3.4.2	Implementation Details	8
4	Results	9
4.1	Overview	9
4.2	Monte Carlo Learning Parameter	9
4.3	Policies	9
4.4	Time	9
4.5	Win/Learning Rate	9
5	Analysis	10
6	Conclusion	11
6.1	Future Work	11
6.2	Personal Reflections	11
7	Acknowledgements	13

Chapter 1

Introduction

Chapter 2

Background

2.1 Reinforcement Learning

2.2 Learning Algorithms

2.2.1 Monte Carlo Learning

Chapter 3

Design & Implementation

This section describes the implementations of Tic-Tac-Toe and Chung Toi as well as the game-playing agents used in this project’s experiments. The source code is not provided in this report, but is available upon request.

3.1 Overview

The project is built in a modular way, using an object-oriented style. All actions, agents, games, and states implement general interfaces, called *Action*, *Agent*, *Game*, and *State* respectively, that serve as Facades for the complex implementations.

The *Game* interface has a method *play()* that simulates a single game. This is where the agents are asked to choose actions, receive returns, and complete other tasks, depending on the specific game. Agents do not distinguish between different games and only use methods defined in the general interfaces, which means that the *MonteCarloAgent* class, for example, can be used to play both Tic-Tac-Toe and Chung Toi. This modularity provides a nice separation of concerns that enables simpler testing and debugging.

3.2 Agents

As part of this project, two agents were implemented: *RandomAgent* and *MonteCarloAgent*.

MonteCarloAgent implements the ϵ -soft on-policy Monte Carlo control algorithm. In its *chooseAction(State s)* method, it chooses an action according to its policy for states it has encountered in previous games, and randomly otherwise. In its *gameOver()* method it calls two private helper methods *policyEvaluation()* and *policyImprovement()* that update Q and π , respectively, as described in section 2.2.1. Hash maps are used to store π and Q for constant time access and insertion.

As its name suggests, *RandomAgent* randomly selects an action from the list of available actions at any given state. In total, the logic for this agent is equivalent to one line of code and this agent is only used for training and benchmarking purposes.

3.3 Tic-Tac-Toe

The first game used in this project is Tic-Tac-Toe. It was chosen because it has a small state space (less than $3^9 = 19,683$ states) and it is a conceptually simple game that most people are familiar with.

3.3.1 Rules

Although this is a widely-known game, the rules are still included here for reference and clarity.

First, one of the two players is randomly selected to use **X** tokens (the other player uses **O** tokens). From now on, these players are referred to as the “X-player” and “O-player”, respectively. The game then consists of the two players alternately placing their tokens in empty spaces on a 3x3 grid, starting with the X-player. Figure 3.1 is an example of a possible state of the grid.

		X
X	O	
O		X

Figure 3.1: Example of a Tic-Tac-Toe grid

The game has three possible endings:

- **X-player wins:** three **X** tokens form a horizontal, vertical, or diagonal row
- **O-player wins:** three **O** tokens form a horizontal, vertical, or diagonal row
- **Draw:** the grid is filled and neither player has won

3.3.2 Implementation Details

The aim of this project is to investigate the effect of eliminating a game's state symmetries on the Monte Carlo algorithm's learning rate and optimal learning parameters. In the game of Tic-Tac-Toe, a state can be symmetrical with up to seven other states as result of flipping it in the following ways:

- horizontal axis
- vertical axis
- horizontal axis and then vertical axis
- major diagonal (i.e. top left to bottom right)
- major diagonal and then horizontal axis
- minor diagonal (i.e. top right to bottom left)
- minor diagonal and then horizontal axis

Of course, for some states, the result of some of these flips are the same, which is why a state might be symmetrical with strictly less than seven other states.

Thus, Tic-Tac-Toe was implemented in three ways:

- **“Normal”**: This version does not break any symmetry and was used as a baseline.
- **“Limited Actions”**: This version breaks the symmetry between states by limiting the actions available from each state such that no two actions result in symmetrical states.
- **“Symmetric Equality”**: This version breaks the symmetry between states by considering two state instances to be equal (i.e. to represent the same state) if they are symmetrical.

An interesting result of the Limited Actions implementation is that it slightly alters the way the game is played. Suppose a *RandomAgent* is asked to select a move from the game's initial state (i.e. empty grid). In the Normal implementation, there are nine possible actions and in the Limited Actions implementation, there are only three possible actions: putting an **X** token in a corner, in the middle of the grid, and in the middle of one of the sides. To give a specific example of how this changes the game, notice that the probability of the agent choosing the action that involves placing their **X** token in the middle of the grid is $\frac{1}{9}$ in the Normal implementation and $\frac{1}{3}$ in the Limited Actions implementation. Thus, the agent will end up choosing this action three times as frequently in the latter case. However,

this effect is ignored for the purposes of this project because in the limit of infinitely many games, the *MonteCarloAgent* will encounter all states and choose all possible moves an equal number of times, thus computing an accurate expected return for each action.

Implementing the desired behaviour for the Symmetric Equality version was not as simple as initially expected. Overriding the *equals()* method inherited from Java’s *Object* class was straightforward. However, overriding the inherited *hashCode()* method was tricky. A Java *HashMap* considers two keys to be distinct if their *hashCode()* methods produce different results or if they produce the same result and their *equals()* methods produce different results. Thus, it was crucial for the *hashCode()* method to produce exactly the same result for symmetrical states. This was done by internally converting each state to a canonical form in such a way that symmetrical states have the same canonical form. Thus,

The idea behind the Limited Actions implementation is that it is expected to speed up the *MonteCarloAgent*’s learning since only a fraction of the states will be reachable. Thus, after a fixed number of games, each possible state will be encountered more often than in the Normal implementation and so the expected return for each state will be more accurate. On the other hand, the Symmetric Equality implementation is expected speed up learning by allowing the *MonteCarloAgent* to combine its learning from symmetric states therefore increasing the accuracy of its expected returns. In both cases, the agent’s policy and action-value function will be smaller, which reduces the agent’s memory requirements.

3.4 Chung Toi

The second game used in the project is Chung Toi, a more complex version of Tic-Tac-Toe that has been studied in the context of reinforcement learning before. Its state space is larger than that of Tic-Tac-Toe (i.e. less than $5^9 = 1,953,125$ states) which leads to interesting results.

3.4.1 Rules

The main idea behind Chung Toi is the same as that behind Tic-Tac-Toe in the sense that the game is played on a 3x3 grid and each player’s goal is to get three of their pieces in a row. However, that is the entire extent of their overlap and although they may seem similar, finding a good strategy for Chung Toi is not at all intuitive.

The game is played as follows: first, one of the two players is randomly selected to use red tokens (the other player uses white tokens). From now on, these players are referred to as the “red player” and “white player”, respectively. Each player has three tokens at their disposal and each token can be used in two possible orientations, as shown in Figure 3.2.




Figure 3.2: Red and white Chung Toi tokens in both possible orientations

The game then consists of two phases:

- Phase 1:
- Phase 2:

Figure 3.3 is an example of a possible state of the grid.




Figure 3.3: Example of a Chung Toi grid

3.4.2 Implementation Details

Chapter 4

Results

4.1 Overview

4.2 Monte Carlo Learning Parameter

4.3 Policies

4.4 Time

4.5 Win/Learning Rate

Chapter 5

Analysis

Chapter 6

Conclusion

6.1 Future Work

6.2 Personal Reflections

This has been a really great learning experience, not only in terms of discovering the field of reinforcement learning and its applications, but also in terms of gaining experience with building up a large, complex programming project from scratch using a variety of technical tools.

When I was in high school, I designed and implemented the classic board game Nine Mens Morris where a user plays against an algorithm I wrote. My plan for the project was to ask my friends and family to play my game and to store the moves made in each game so that I could build up a database of sequences of moves and outcomes of the game. I hoped to somehow compute which moves were the best and get my algorithm to learn which moves made winning the most likely. Little did I know, I wanted to invent RL from scratch unaware of its existence or of the amount of research in the field. Several years later, I am delighted to have had the chance to work on a similar project, but this time on a deeper level using well-known computational algorithms.

When I started working on this project, I decided to treat it as an opportunity to develop many skills that are critical for a career in software engineering:

- Writing clean, maintainable, well-documented code
- Designing and implementing tests for my code
- Using a build tool to manage dependencies between packages in my project as well as with external libraries
- Using version control effectively

I tested the methods in my API using a unit-testing framework for Java called JUnit. I also used a testing framework for Java called Mockito this allowed me to verify the behaviour of objects with external dependencies by creating mock objects for these dependencies, which mimic real objects but do so in a particular way that I can specify.

In order to save my experiment results in a format that would facilitate the creation of graphs, I used opencsv, a CSV parser library for Java.

To manage my projects dependencies, I used a build tool developed by Google called Bazel and I used Git for version control.

Overall, I really enjoyed learning about RL algorithms and exploring their applications and success rates, all while developing strong programming skills that will help me throughout my career.

Chapter 7

Acknowledgements

Thanks for reading.